

CSE 489/589 Spring 2014
Programming Assignment 3
Routing Protocols

Due Time: 05/02/2014 @ 23:59:59

1. Problem Statement

In this assignment you will implement a simplified version of the *Distance Vector Protocol*. The protocol will be run on top of servers (behaving as routers) using UDP. Each server runs on a machine at a pre-defined port number. The servers should be able to output their forwarding tables along with the cost and should be robust to link changes. (Note: we would like you to implement the basic algorithm: count to infinity **not** poison reverse. In addition, a server should send out routing packets only in the following two conditions: **a) periodic update** and **b) the user uses command asking for one**. This is a little different from the original algorithm which immediately sends out update routing information when routing table changes.

NOTE: Use UDP Sockets only for your implementation. You should **use only the select() API** for handling multiple socket connections. Please don't use multi-threading or fork-exec.

If you're using C++, you are not allowed to use any STLs (Standard Template Library) for the socket programming part. **If you use any STLs for the socket programming part, your project will not be graded.**

2. Getting Started

A Distance Vector Routing Algorithm

Text book: Page 371 – Page 377.

3. Protocol Specification

The various components of the protocol are explained step by step. Please strictly adhere to the specifications.

3.1 Topology Establishment

In this programming assignment, you will use **five** CSE student servers – {**timberlake, nickelback, metallica, dragonforce, beatles**}.cse.buffalo.edu. Each server is supplied with a topology file at startup that it uses to build its initial routing table. The topology file is local and contains the link cost to the neighbors. For all other servers in the network, the initial cost would be infinity. Each server can only read the topology file for itself. The entries of a topology file are listed below:

- **<num-servers>**
- **<num-neighbors>**
- **<server-ID> <server-IP> <server-port>**

- `<server-ID1> <server-ID2> <cost>`

num-servers: total number of servers.

server-ID, server-ID1, server-ID2: a unique identifier for a server, which is assigned by you.

cost: cost of a given link between a pair of servers. Assume that cost is an integer value.

E.g., consider the topology in Figure 1. We give a topology file for server 1(timberlake).




Figure 1: Example topology

| Line number | Line entry | Comments |
|-------------|----------------------|---|
| 1 | 5 | number of servers |
| 2 | 3 | number of edges or neighbors |
| 3 | 1 128.205.36.8 4091 | server-id 1 and corresponding IP, port pair |
| 4 | 2 128.205.35.24 4094 | server-id 2 and corresponding IP, port pair |
| 5 | 3 128.205.36.24 4096 | server-id 3 and corresponding IP, port pair |
| 6 | 4 128.205.36.4 7091 | server-id 4 and corresponding IP, port pair |
| 7 | 5 128.205.36.25 7864 | server-id 5 and corresponding IP, port pair |
| 8 | 1 2 7 | server-id and neighbor id and cost |
| 9 | 1 3 4 | server-id and neighbor id and cost |
| 10 | 1 4 5 | server-id and neighbor and cost |

Your topology files should only contain the Line entry part (2nd column, see topology_example.txt). In each line, every two elements (e.g., server-id and corresponding IP, corresponding IP and port number) should be separated with a **space**. For cost values, **each topology file should only contain the cost values of the host server's neighbors** (The host server here is the one which will read this topology file). You can use your own topology files to test your code. However, we will use our topology files to test your program. So please adhere to the format for your topology files.

IMPORTANT: In this environment, costs are bi-directional i.e. the cost of a link from A-B is the same for B-A. Whenever a new server is added to the network, it will read its topology file to determine who its neighbors are. Routing updates are exchanged periodically between neighboring servers. When this newly added server sends routing messages to its neighbors, they will add an entry in their routing tables corresponding to it. Servers can also be removed from a network. When a server has been removed from a network, it will no longer send distance vector updates to its neighbors. When a server no longer receives distance vector updates from its neighbor for three consecutive update intervals, it assumes that the neighbor no longer exists in the network and makes the appropriate changes to its routing table (link cost to this neighbor will now be set to infinity but not remove it from the table). This information is propagated to other servers in the network with the exchange of routing updates. Please note that although a server might be specified as a neighbor with a valid link cost in the topology file, the absence of three consecutive routing updates from this server will imply that it is no longer present in the network.

3.2 Routing Update

Routing updates are exchanged periodically between neighboring servers based on a time interval specified at the startup. In addition to exchanging distance vector updates, servers must also be able to respond to user-specified events. There are 4 possible events in this system. They can be grouped into three classes: topology changes, queries, and exchange commands. Topology changes refer to an updating of link status (update). Queries include the ability to ask a server for its current routing table (display), and to ask a server for the number of distance vectors it has received (packets). In the case of the packets command, the value is reset to **zero** by a server after it satisfies the query. Exchange commands can cause a server to send distance vectors to its neighbors immediately. Examples of these commands include:

- **update 1 2 inf**
The link between the servers with IDs 1 and 2 is assigned to infinity.
- **update 1 2 8**
Change the cost of the link to 8.
- **step**
Send routing update to neighbors right away. Note that except this, routing updates only happen periodically.
- **packets**
Display the number of distance vector packets this server has received since the last instance when this information was requested.
- **disable server-id**
Disable the link to given server. Here you need to check if the given server is its neighbor.
- **crash**
Emulate a server crash. Close all connections on all links. The neighboring servers must handle this close correctly and set the link cost to infinity.

- **display**

Display the current routing table. The display should be formatted as a sequence of lines, with each line indicating: <destination-server-ID> <next-hop-server-ID> <cost-of-path>

3.3 Message Format

Routing updates are sent using the General Message format. All routing updates are UDP unreliable messages. The message format for the data part is:

0 1 2 3 (10 bits)
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 (bit)

| | |
|-------------------------|-------------|
| Number of update fields | Server port |
| Server IP | |
| Server IP address 1 | |
| Server port 1 | 0x0 |
| Server ID 1 | Cost 1 |
| Server IP address 2 | |
| Server port 2 | 0x0 |
| Server ID 2 | Cost 2 |
| | |

- **Number of update fields:** (2 bytes): Indicate the number of entries that follow.
- **Server port:** (2 bytes) port of the server sending this packet.
- **Server IP:** (4 bytes) IP of the server sending this packet.
- **Server IP address n:** (4 bytes) IP of the n-th server in its routing table.
- **Server port n:** (2 bytes) port of the n-th server in its routing table.
- **Server ID n:** (2 bytes) server id of the n-th server on the network.
- **Cost n:** cost of the **path** from the server sending the update to the n-th server whose ID is given in the packet.

Note: First, the servers listed in the packet can be any order i.e., 5, 3, 2, 1, 4. Second, the packet needs to include an entry to reach itself with cost 0 i.e. server 1 needs to have an entry of cost 0 to reach server 1.

4. Server Commands/Input Format

To start a server, type:

- ***.server -t <topology-file-name> -i <routing-update-interval>***

topology-file-name: The topology file contains the initial topology configuration for the server, e.g., timberlake_init.txt. Please adhere to the format described in 3.1 for your topology files.

routing-update-interval: It specifies the time interval between routing updates in seconds.
port and server-id: They are written in the topology file. The server should find its port and server-id in the topology file without changing the entry format or adding any new entries.

The following commands can be specified at any point during the run of the server:

- **update** *<server-ID1> <server-ID2> <Link Cost>*
server-ID1, server-ID2: The link for which the cost is being updated.
Link Cost: It specifies the new link cost between the source and the destination server. Note that this command will be issued to **both** *server-ID1* and *server-ID2* and involve them to update the cost and no other server.
- **step**
Send routing update to neighbors (triggered/force update)
- **packets**
Display the number of distance vector packets this server has received since the last invocation of this information.
- **display**
Display the current routing table. And the table should be displayed in a **sorted** order from small ID to big.
- **disable** *<server-ID>*
Disable the link to a given server. Doing this “closes” the connection to a given server with *server-ID*.
- **crash**
“Close” all connections. This is to simulate server crashes.

5. Server Responses/Output Format

The following are a list of possible responses a user can receive from a server:

- On successful execution of an update, step, packets, display or disable command, the server must display the following message:

***<command-string>* SUCCESS**

where *command-string* is the command executed. Additional output as desired (e.g., for display, packets, etc. commands) is specified in the previous section.

- Upon encountering an error during execution of one of these commands, the server must display the following response:

<command-string> <error message>

where *error message* is a brief description of the error encountered.

- On successfully receiving a route update message from neighbors, the server must display the following response:

RECEIVED A MESSAGE FROM SERVER *<server-ID>*

where the *server-ID* is the id of the server which sent a route update message to the local server.

6. Submission

- You have to submit a working implementation of the project.
- You also have to submit a short report explaining the overall implementation details. Your report should mention the file name and the line number where you define the data structure of the update message and the data structure of the routing table.
- Name your main file as *<ubit_name>_server.c* or *<ubit_name>_server.cpp*
- Name your executable as '*server*'
- Name your report as *<ubit_name>_report.pdf*
- Makefile is compulsory. If there is no makefile, you'll lose 15% points.
- You should combine all your submission materials into one '.tar' file. Name your .tar file as *<ubit_name>_proj3.tar*. Use the submission command, `submit_cse489` or `submit_cse589`, to submit your tar file.
- You DO NOT have to submit any topology files.